

01 | THE PRODUCTION FUNCTION IS CHANGING

The economics of software delivery are being fundamentally restructured — not incrementally improved.

THE OLD FORMULA

$$\text{Human Time} \times \text{Hourly Rate} = \text{Delivery Cost}$$



THE NEW FORMULA

$$\begin{aligned} & (\text{Human Orchestration} \times \text{Strategic Rate}) \\ & + (\text{Agent Development} \times \text{Amortisation}) \\ & + (\text{Agent Run} \times \text{Consumption Rate}) \\ & + (\text{Human Validation} \times \text{Review Rate}) \end{aligned}$$



Calibration Baselines Drifting

Velocity data calibrated to human effort becomes invalid when agents execute 40–60% of implementation tasks.



Cost Structure Inverted

AI-intensive projects spend more on tooling & agents and less on human labour — breaking every budget template you have.



Non-Linear Cost Curves

Multi-agent chains create exponential cost relationships. A small change in task complexity produces a large, unpredictable cost spike.



Pricing Model Obsolescence

Clients will not pay time & materials for tasks agents execute. Outcome-based contracts are accelerating in the market.

02 | THE THREE-LAYER DELIVERY STACK

All three layers coexist permanently. The mix shifts — they do not replace each other.

L1



HUMAN STRATEGIC LAYER

Who: Senior engineers, architects, product leads, delivery managers

Does: Judgment • Creativity • Stakeholder relationships • Architecture decisions • Agent orchestration • Outcome ownership

Economics: Day rate / outcome fee. Highest cost per person but irreducible.

Delivery Mix %

2026: 35%

2028: 25%

2030: 20%

L2



AI-AUGMENTED HUMAN LAYER

Who: Developers, QA engineers, analysts working with AI tools

Does: Code generation assist • Test automation • Documentation • Pattern matching • Scaffolding

Economics: Day rate + AI tool consumption cost. 2–5x human throughput for suitable tasks.

Delivery Mix %

2026: 50%

2028: 40%

2030: 30%

L3



AUTONOMOUS AGENT LAYER

Who: AI agents operating independently, spawning sub-agents, interacting with systems

Does: Autonomous coding • PR creation • Test execution • System integration • Validation pipelines

Economics: Per-task / per-compute consumption. Variable, non-linear, FinOps-managed.

Delivery Mix %

2026: 5%

2028: 25%

2030: 40%

KEY INSIGHT: By 2030, agent-executed work (L3) may account for 40%+ of delivery activity on AI-native projects — yet most estimation models today have zero visibility into this layer.

03 | TOMORROW'S PRICING MODELS — Three Paradigms Emerging Simultaneously



HUMAN PRICING

From Time → Outcomes

Time & Materials

Transitioning out — agents make human time-billing untenable for implementation tasks

Fixed Price

Provider carries delivery risk — requires robust probabilistic estimation


Outcome-Based

Payment tied to measurable business result — fastest growing model

Value-Based

Priced against business value delivered — requires outcome measurement capability

Estimation shift: From effort prediction → outcome probability modelling



AI TOOL PRICING

From Seats → Consumption

Today — Seat-Based

~£15–20/user/month fixed. Treated as overhead. No project-level signal.

Tomorrow — Consumption

Variable cost tied to tokens, tasks, compute. Becomes a project variable.

High-AI Projects

More tooling cost, less human cost. Tradeoff must be modelled explicitly.

Estimation Impact

AI tool cost moves from fixed overhead to project-level variable — must be tracked and forecast.

Estimation shift: AI tooling becomes a project cost line, not an overhead



AGENT PRICING

Three Models Converging

Per-Task

Fixed fee per agent task completed. Requires well-scoped tasks. Predictable if taxonomy is robust.

Per-Compute

Tokens + API calls + memory. Current API model. Non-linear with chain depth.

Per-Outcome

Pay only on verified result. Logical endpoint — eliminates cost risk for buyers.

The Risk

Multi-agent chains multiply costs non-linearly. A 'simple' task can cost \$0.10 or \$50+.

Estimation shift: Agent cost needs FinOps disciplines — not effort estimation

04 | THE CRITICAL DISTINCTION — Agent Development vs. Agent Run Cost

These are categorically different economic events requiring different modelling, governance, and organisational ownership.



AGENT DEVELOPMENT COST

Capital Investment — Amortised over task volume

Prompt Engineering & System Design

Human effort — senior engineer hours. High uncertainty on first-of-type.

Tool & API Integration

Standard integration effort. Better understood than prompt engineering.

Guardrail & Safety Configuration

Compliance-heavy environments add 30–50% to baseline.

Testing & Validation

40–60% of total build cost. Significantly more than traditional testing.

Orchestration Architecture

Highest uncertainty. Novel design — reference class data is scarce.



AGENT RUN COST

Operating Expenditure — Managed like cloud infrastructure

Predict.

Token Consumption

Context window depth × reasoning chain length. Moderate predictability.

Low–Med

Tool Call Frequency

External API calls per task. Low on novel tasks, improves with calibration.

Low

Sub-Agent Spawning

Non-linear cost. Depth limits and cost guardrails essential.

Very Low

Failure & Retry Cost

Agent reliability rate × error handling cost. Improves with familiarity.

Very Low

Memory & State Mgmt

Context persistence for long-running agents. Requires

Low–Med

$$\text{AGENT ROI} = (\text{Human cost per task} \times \text{projected volume}) - (\text{Development cost} + \text{Run cost} \times \text{projected volume}) \quad | \quad \text{Treat development as CAPEX, run as OPEX}$$

05 | THE FOUR-DIMENSION COST VECTOR — The New Estimation Unit


A single effort number is no longer sufficient. Every project needs a four-dimension cost vector.

D1 Human Orchestration Senior human time for strategy, architecture, stakeholder management, and agent oversight Pricing Basis: Day rate / outcome fee Predictability: High — most stable Governance Owner: Delivery Lead Now: 35–40% 2030: 20–30%	D2 Agent Development One-time investment to build and validate agent capability. Amortised across task volume. Pricing Basis: Human effort at build cost Predictability: Moderate — high first-of-type uncertainty Governance Owner: Engineering Lead Now: 5–15% 2030: 10–15%	D3 Agent Run Ongoing variable consumption cost of operating deployed agents at scale. Pricing Basis: Compute / token / task consumption Predictability: Low initially — FinOps-managed Governance Owner: FinOps Team Now: 2–8% 2030: 30–45%	D4 Human Validation Human review, testing, and governance of agent outputs. Scales with agent output volume. Pricing Basis: Day rate Predictability: Moderate — scales with reliability Governance Owner: QA / Delivery Lead Now: 45–55% 2030: 15–25%
--	---	---	---

By 2029, D3 Agent Run Cost may be the single largest variable cost on AI-intensive projects — larger than human implementation labour. Most estimation functions today have zero visibility into this dimension.

INSIGHT | HOW PROJECT COST MIX WILL SHIFT — 2026 to 2030

The relative weight of each cost dimension shifts dramatically over 4 years. Estimation models built today will be systematically wrong by 2028.



WATCH D3: Agent Run Cost grows from 5% to 40% of project cost by 2030 — the fastest-growing, least-understood, and currently unmodelled cost dimension.

RISK | FIVE NEW ESTIMATION RISKS THE AGENT ERA INTRODUCES

Current estimation frameworks cannot accommodate these risks. They require structural model changes.



Agent Reliability Risk

Agents fail, hallucinate, and produce outputs requiring human intervention. Failure rates are task- and context-dependent and hard to predict upfront.

Mitigation: Include failure and retry cost buffer in all agent-intensive estimates. Calibrate from deployment actuals.



Task Decomposability Risk

Cost and feasibility of automation depends entirely on whether a task can be decomposed into well-scoped, verifiable sub-tasks — often discovered late.

Mitigation: Mandatory agent suitability assessment at intake. Tier 3 and 4 items flagged explicitly before estimation begins.



Cost Non-Linearity Risk

Multi-agent chains create exponential cost relationships. Small changes in task complexity produce large, unpredictable changes in agent run cost.

Mitigation: Scenario-based run cost modelling. Explicit cost ceiling guardrails per agent invocation at pipeline design.



Context Bleed Cost

Long-running agents accumulate context requiring expensive large context windows or complex memory management — an invisible but real cost.

Mitigation: Context architecture cost included in agent development estimate. Memory management modelled as ongoing run cost.



Governance & Compliance Overhead

Agents taking autonomous actions create audit, accountability, and compliance requirements — material cost in regulated industries, currently invisible.

Mitigation: Separate governance cost estimate based on regulatory context and agent action scope. Do not fold into development estimate.

07 | ESTIMATION FUNCTION — WHAT TO BUILD AND WHEN

Three capability horizons. Start immediately — calibration data must be collected from real deployments before scale makes the problem unmanageable.

HORIZON 1

2026 — Foundation

Agent Cost Awareness

Educate estimation team on dev vs. run distinction and four-dimension cost vector — critical immediately

Agent Task Taxonomy v1.0

Design and pilot four-tier suitability framework. Integrate into project intake for all AI components

Consumption Instrumentation

Instrument any existing AI usage to capture per-project consumption data — the raw material for calibration

FinOps Connection

Establish FinOps partnership. Align cost taxonomy before agent deployment scales — before, not after

First Agent ROI Model

Build ROI projection for any active agent investment. Use as template for all future decisions

HORIZON 2

2027 — Calibration

Agent Run Cost Baselines

From 6+ months of instrumented usage, build cost-per-task baselines by taxonomy tier

Failure Rate Calibration

Measure agent failure and escalation rates by task type. Incorporate cost buffers into estimation models

Four-Dimension Cost Vector

Deploy on pilot projects. Compare forecast to actuals. Calibrate dimension weights empirically

Scenario-Based Run Models

Develop low/base/high scenarios for run cost — genuine uncertainty requires ranges, not points

Outcome Probability Modelling

Begin statistical capability to forecast outcome probability — precursor to outcome-based pricing

HORIZON 3

2028–30 — Intelligence

Unified Delivery Cost Model

Full four-dimension cost vector with calibrated probability distributions, updated continuously from actuals

Agent Portfolio Optimisation

Analysis of optimal human/agent mix across portfolio — informs workforce and technology investment strategy

Outcome-Based Contract Pricing

Capability to price and risk-manage outcome-based contracts with delivery probability modelling

AI Leverage Intelligence

Continuous portfolio reporting on AI leverage by layer and task type — drives investment prioritisation

FinOps Convergence

Full integration of delivery estimation and agent cost management into unified financial planning

IMPACT | THE NEW FINANCE CONVERSATION — What Changes for COO & CIO

The estimation function's output to finance changes fundamentally. This requires proactive governance adaptation.

TODAY'S CONVERSATION

"This project requires 8 developers for 6 months at £650/day.
Total: £624,000.
Delivery: Q3 2026."

Finance books a fixed labour cost.
Variance is explained by scope change or delay.

*Governance model: Fixed date commitments.
Budget = headcount × duration.*



TOMORROW'S CONVERSATION

D1 Human Orchestration:	£180,000 (fixed)
D2 Agent Development:	£95,000 (amortised, 18 months)
D3 Agent Run Cost:	£45,000–£120,000 (variable)
D4 Human Validation:	£60,000 (semi-fixed)

P50 Total: £330,000 | P85 Total: £455,000–£530,000

Finance must now model:

- Variable D3 costs with scenario ranges
- CAPEX/OPEX split for agent investment
- Probability-range project budgets

SUMMARY | THE AGENT-ERA ESTIMATION TRANSFORMATION AT A GLANCE

Dimension	Current State	Agent-Era End State	Key Action	Horizon
Delivery Model	Human effort dominant. AI as productivity assist.	Three-layer stack: human strategic, AI-augmented, autonomous agent.	Build three-layer delivery taxonomy and resourcing model.	2026
Pricing Model	Time & materials. Seat-based AI tooling.	Outcome-based human pricing. Consumption-based agent pricing.	Develop outcome probability modelling capability.	2027–28
Estimation Unit	Single effort number. Story points or hours.	Four-dimension cost vector: D1 + D2 + D3 + D4.	Restructure estimation model around four dimensions.	2026–27
Agent Economics	Not modelled. Dev and run cost invisible.	Agent ROI model with amortised development + variable run cost.	Build ROI framework. Instrument consumption now.	2026
Cost Predictability	High for human labour. AI tooling as fixed overhead.	High D1/D4. Improving D2. Low but managed D3.	FinOps integration. Scenario-based run cost modelling.	2026–27
Finance Interface	Fixed-date, headcount-based project budgets.	Probability ranges + CAPEX/OPEX split + variable D3 budgets.	Finance working group. Governance adaptation.	2026
Function Mandate	Project estimation and portfolio tracking.	Strategic intelligence: delivery, AI leverage, agent economics, contract pricing.	Expand scope and skill base in parallel with model transformation.	2027–28

The strategic window is open now. Calibration data must be collected from real deployments — the sooner measurement begins, the sooner the model becomes reliable.

LAYER 1 | MEASUREMENT INFRASTRUCTURE — Instrument Before You Estimate

You cannot estimate what you cannot measure. Deploy telemetry first — every other layer depends on this data.

WHAT EVERY AGENT PIPELINE MUST EMIT

Input & Output Tokens:	By model — primary and secondary cost drivers
Tool Calls:	Count and type — often priced separately
Sub-Agent Spawns:	Recursively tracked — multiplier on all costs
Retry Count:	Failures and retries — most underestimated component
Wall-Clock Time:	For infrastructure cost and timeout modelling
Task ID & Taxonomy Tier:	Enables per-tier calibration of reference table
Project Attribution:	Required for FinOps chargeback
Model Used:	Different models = different unit costs
Success / Failure Flag:	Reliability measurement and retry cost modelling
Total Calculated Cost:	Primary output — attributed to project and task type

MINIMUM VIABLE DASHBOARD

Total tokens consumed by model & project

Tool calls by type & project

Cost per invocation & per task completion

Failure rate and retry cost by task type

7-day cost trend per project

DATA QUALITY WARNING

If telemetry does not distinguish actual task completion from invocation count, all cost calculations are distorted. Validate a sample of 50+ invocations manually before building any model. Poor instrumentation is worse than no instrumentation.

IMPLEMENTATION: Months 1–3. Do not attempt to estimate during this phase — only collect, validate, and observe.

LAYER 2 | TASK TAXONOMY & COST CLASSIFICATION — Classify Before You Model

Every agent task must be assessed across five cost-relevant dimensions before estimation begins. The classification drives everything downstream.

A

Context Volume

Levels: Small (<10K) → Medium → Large → Very Large (200K+)

Most controllable cost driver. Very Large = mandatory architecture review.

B

Reasoning Depth

Levels: Shallow (1–3 steps, 1×) → Moderate (1.5–3×) → Deep (3–8×) → Unbounded (ceiling req'd)

Drives output token count. Deep + unbounded = mandatory cost ceiling.

C

Tool Call Intensity

Levels: Minimal (0–2) → Moderate (3–10) → Heavy (10–50) → Very Heavy (50+ or recursive)

Failure probability rises with call count. Very Heavy = architecture review first.

D

Sub-Agent Dependency

Levels: None → Fixed (1–3 agents) → Dynamic (agent decides spawn count)

Creates recursive cost structures. Dynamic spawn = hard depth limit mandatory.

E

Iteration Likelihood

Levels: Low (10–15% buffer) → Medium (20–40%) → High (50–100%) → Very High (don't model as agent)

The most underestimated component. Must include retry cost buffer explicitly.

 TASK CLASSIFIER — Classify tasks across 5 dimensions to assign Cost Class and estimation

DIMENSION REFERENCE — Use dropdowns in the Classifier table below

A: Context Volume	Small (<10K) → Negligible cost	Medium (10K-50K) → Moderate cost	Large (50K-200K) → Significant — compression recommended	Very Large (200K+) → Architecture review mandatory	
B: Reasoning Depth	Shallow (1-3 steps) → 1.0x multiplier	Moderate (4-10 steps) → 1.5-3.0x multiplier	Deep (10+ steps) → 3-8x multiplier	Unbounded → Hard ceiling required	
C: Tool Call Intensity	Minimal (0-2) → Negligible	Moderate (3-10) → Budget in cost model	Heavy (10-50) → Scenario model required	Very Heavy (50+) → Architecture review first	
D: Sub-Agent Dependency	None → No multiplication	Fixed (1-3 agents) → Model each sub-agent	Dynamic (agent decides) → Hard depth limit mandatory		
E: Iteration Likelihood	Low (well-defined) → +10-15% buffer	Medium (partial def) → +20-40% buffer	High (novel/ambiguous) → +50-100% buffer	Very High → Do not model as agent	

 TASK CLASSIFIER — Classify tasks across 5 dimensions to assign Cost Class and estimation

DIMENSION REFERENCE — Use dropdowns in the Classifier table below

A: Context Volume	Small (<10K) → Negligible cost	Medium (10K-50K) → Moderate cost	Large (50K-200K) → Significant — compression recommended	Very Large (200K+) → Architecture review mandatory	
B: Reasoning Depth	Shallow (1-3 steps) → 1.0x multiplier	Moderate (4-10 steps) → 1.5-3.0x multiplier	Deep (10+ steps) → 3-8x multiplier	Unbounded → Hard ceiling required	
C: Tool Call Intensity	Minimal (0-2) → Negligible	Moderate (3-10) → Budget in cost model	Heavy (10-50) → Scenario model required	Very Heavy (50+) → Architecture review first	

PROJECT DETAILS

QUICK SUMMARY

Project Name:		Total P50 (Optimistic):	£0.00
Project Code:		Total P70 (Base Budget):	£0.00
Estimation Date:		Total P90 (Stress / Ceiling):	£0.00
Maturity Stage:	First deployment (0-4 weeks)		
Estimation Lead:			

→ Hard		
+50-	Very High → Do not model as agent	

TASK ROLLUP — Summary from Cost Calculator (manual copy or reference)

Task Name	Cost Class	Volume/Week	Weekly Median	Weekly P70	Weekly P90	4-Week Total (P70)	Notes
		1					

LAYER 3 | COST DRIVER DECOMPOSITION MODEL — The Five-Component Formula

Total Agent Run Cost per Task = C_tokens + C_tools + C_infra + C_retry + C_agents



C_tokens

Token Cost

Formula:

(Input tokens × input price) +
(Output tokens × output price)

Primary driver. Varies by model tier. Route Class 1 tasks to cheaper models (Haiku, GPT-4o-mini) for 70–90% reduction.

Est. Risk: Medium



C_tools

Tool Call Cost

Formula:

Σ (calls of type n × cost per call type n)

Many calls are free but raise failure probability. LLM-as-tool calls multiply token cost — track separately.

Est. Risk: Medium



C_infra

Infrastructure Cost

Formula:

Wall-clock seconds × compute rate per second

Minimal for pure API agents. Significant for self-hosted models or persistent agent runtimes.

Est. Risk: Low



C_retry

Retry & Failure Cost

Formula:

Expected retry rate × tokens per attempt × avg retries to success

THE most underestimated component. 20% failure rate + 2.5 avg retries = +50% on baseline cost. Must be budgeted explicitly.

Est. Risk: High



C_agents

Sub-Agent Cost

Formula:

Σ (sub-agent invocations × full cost per sub-agent)

Recursive: apply all five components at each spawn level. 3-level fan-out can multiply parent cost by 9×. Hard depth limits essential.

Est. Risk: Very High

COMPOUNDING EXAMPLE: Parent (£2) spawns 3 sub-agents (£1.50 each) each spawning 2 workers (£0.80 each) = Grand total £11.30 from a task budgeted at £2.00.

LAYER 3 | REFERENCE COST TABLE — Your Estimation Lookup Tool (Calibrate from Actuals)

Initially populated with reference estimates. Replace with empirical actuals from L1 instrumentation within 6 months. Recalibrate quarterly.

Task Type	Tier	Class	Median Cost	P70 Cost	P90 Cost	Primary Driver	Estimation Method
Simple retrieval / Q&A	1	1	£0.05	£0.09	£0.18	Input tokens	Point estimate — high confidence
Code gen — single function	1	2	£0.45	£0.80	£1.60	Context + output	Range estimate ±30%
Code review — file-level	1	2	£1.20	£2.10	£4.50	Context volume	Range estimate ±30%
Bug fix / debugging	1	3	£3.50	£7.00	£18.00	Retry cost	Scenario model — high retry variance
Feature implementation	2	3	£12.00	£25.00	£55.00	Context + retries	Scenario model — document assumptions
Architecture analysis	1	3	£18.00	£38.00	£85.00	Context (very large)	Scenario model — compression mandatory
Multi-file refactoring	2	4	£45.00	£100.00	£250.00	Context + sub-agents	Cost ceiling + continuous monitoring
Test suite generation	2	3	£8.00	£18.00	£40.00	Output tokens + tools	Scenario model
CI/CD pipeline config	1	1	£0.80	£1.40	£2.80	Tokens (predictable)	Point estimate
Documentation generation	1	2	£2.50	£4.50	£9.00	Output tokens	Range estimate ±30%

These figures are illustrative starting points. Never use vendor benchmarks as your cost profile — they will be wrong for your architecture. Build from your own L1 actuals.

LAYER 4 | SCENARIO MODELLING & GUARDRAILS — Ranges, Ceilings, Controls

Produce three scenarios — not a point estimate. Implement cost ceilings and guardrails in code, not just in governance documents.

OPTIMISTIC P50

Use for: Internal planning target. Team goals.

Calc: Median cost × volume × 1.0× reliability

BASE P70

Use for: Project budget line. Finance planning.

Calc: P70 cost × volume × 1.3× reliability

STRESS P90

Use for: Cost ceiling basis. Risk reserve. Contracts.

Calc: P90 cost × volume × 1.6–2.0× reliability

THREE-LEVEL COST CEILINGS — Mandatory for Class 4 Tasks and Any Dynamic Sub-Agent Spawning



Per-Invocation:

P90 × 1.5 for that task type. Agent halts and escalates when reached.



Per-Project Daily:

Weekly P70 budget ÷ 5 × 1.5–2.0 buffer. Alert at 70% and 90%.



Per-Project Total:

Stress (P90) scenario total. Finance review triggered when approached.

GUARDRAILS MUST BE IN CODE — not in governance documents. Every threshold must be a programmatic check in the agent orchestration layer.

LAYER 4 | THE FULL GUARDRAIL STACK — Eight Controls, All Required in Code

A guardrail that exists only in a policy document is a wish, not a control. Every item below must be implemented programmatically in the agent orchestration layer.



Token Budget per Invocation

Threshold: P95 estimate + 20% → Agent terminates. Logs partial output. Flags for review.



Tool Call Budget

Threshold: Taxonomy estimate + 50% → Agent pauses. Requests human instruction to continue.



Spawn Depth Limit

Threshold: Never more than 3–4 levels → Sub-agent creation blocked. Parent agent notified.



Retry Limit

Threshold: 3–5 attempts by class → Task marked failed. Human review triggered. Cost logged.



Daily Spend Alert (Warning)

Threshold: 70% of daily ceiling → Email alert to delivery lead. No action required.



Daily Spend Alert (Critical)

Threshold: 90% of daily ceiling → Immediate alert. Agent throttling recommended.



Anomaly Detection

Threshold: 2× 7-day rolling average → Immediate review. Possible prompt issue or model change.



Model Tier Enforcement

Threshold: Class 1 → cheap model only → Alert when expensive model used for Class 1 task.

AUDIT REQUIREMENT: The estimation function should verify programmatic guardrail implementation as part of all project governance gates before agent go-live.

LAYER 5 | PORTFOLIO COST INTELLIGENCE — FinOps Reporting & Optimisation

Four reporting levels connect operational telemetry to strategic decision-making.

Level 1 Invocation	Level 2 Project	Level 3 Portfolio	Level 4 Strategic
Cadence: Real-time Owner: Platform Eng Output: Raw telemetry. Basis for all levels.	Cadence: Weekly Owner: Delivery Lead + Estimation Output: Cost health report. Variance. Forecast update.	Cadence: Monthly Owner: Estimation + FinOps Output: Portfolio dashboard. Optimisation register.	Cadence: Quarterly Owner: Estimation Head + CIO/COO Output: AI leverage intelligence. Investment recommendations.


OPTIMISATION LEVERS — Run Systematically from L1 Data

Model Tier Routing 70–90% Class 1 tasks	Prompt Compression 20–40% After baseline	Context Caching 30–70% Repeated context	Result Caching Variable High-volume T1	Retry Redesign 30–50% After failure data	Spawn Reduction Significant Class 4 tasks	Batching 15–30% Frequent T1 tasks	Context Windowing 20–50% Large context class
--	---	--	---	---	--	--	---

START WITH MODEL TIER ROUTING: Route all Class 1 tasks to cheaper models immediately. Zero architectural risk. 70–90% cost reduction on eligible tasks. No waiting for calibration data.

ROADMAP | IMPLEMENTATION SEQUENCE — Three Phases, 12 Months

Build layers in sequence. Do not skip. Each phase delivers value before the next begins.



THE MOST IMPORTANT STEP: Start Phase 1 this week. Every week of uninstrumented agent workloads is calibration data permanently lost.

SUMMARY | THE AGENT RUN COST ESTIMATION FRAMEWORK AT A GLANCE

Layer	Primary Output	Key Action	Timeline
Layer 1: Measure	Per-invocation cost telemetry by project and task type	Instrument all agent pipelines this month	Months 1–3
Layer 2: Classify	Five-dimension cost class for every agent task	Design taxonomy. Integrate into project intake.	Months 2–5
Layer 3: Decompose	Component estimates: C_tokens + C_tools + C_infra + C_retry + C_agents	Build reference cost table from L1 actuals.	Months 4–6
Layer 4: Control	Three-scenario estimate + cost ceilings + guardrail stack in code	Scenario models + all guardrails implemented in code	Months 5–7
Layer 5: Optimise	Monthly portfolio dashboard + quarterly strategic intelligence	FinOps integration + optimisation lever programme	Months 6–12

QUICK-START ACTIONS — THIS WEEK

1. Pick one live agent workload and instrument it end-to-end
2. Classify that workload's tasks across the five cost dimensions
3. Contact FinOps team — align on attribution taxonomy
4. Set a per-invocation cost ceiling based on observed P95
5. Schedule the Layer 1 dashboard build for next sprint

The framework builds incrementally. The first dataset from one instrumented workload will teach you more than any benchmark. Start with Layer 1 this week.